



04 – Regression

Data Science and Management

Corso di Laurea Magistrale in
Ingegneria Gestionale

Marco Mamei, Natalia
Hadjidimitriou, Fabio
D'Andreagiovanni, Matteo Martinelli

{marco.mamei, selini,
fabio.dandreagiovanni,
matteo.martinelli}@unimore.it

- Regression
- Forecasting

Regression



We speak of **regression** when the target variable to be predicted/forecast is a **real-valued variable**

Examples:

- Predict salary as a function of employee features
- Forecast stock price given historical data

Regression



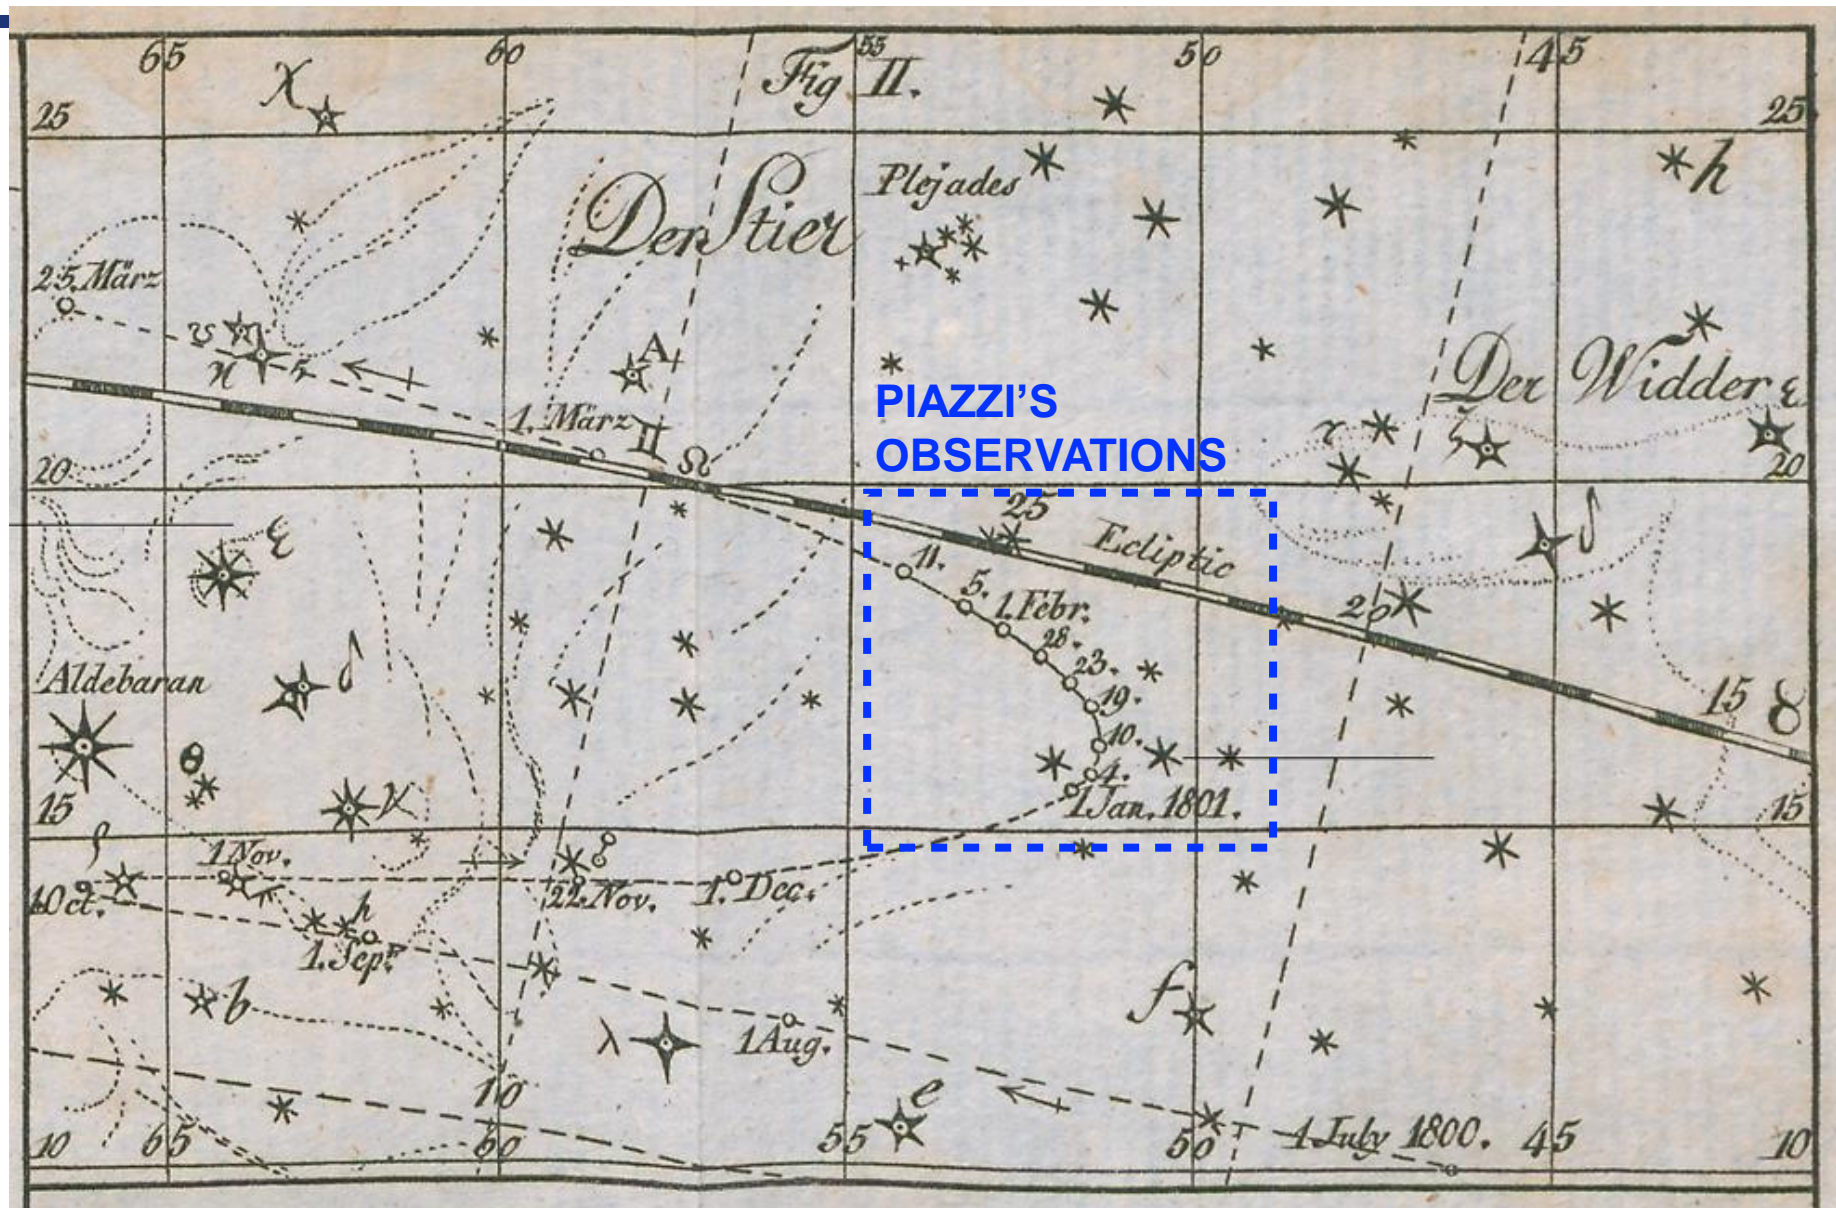
The astronomer **Giuseppe Piazzi** in 1801 discovered what he thought to be a new planet.

He observed the new celestial body for over a month, then it was **lost** in the glare of the Sun.

Gauss computed the orbit of the planet with the **least-squares method** (also developed by Legendre)

That “planet” is now known as the asteroid **Ceres**.

Regression



Regression



Classic approaches

- **KNN-regression**
- **Least-squares linear regression**
- Neural networks
- Regression trees
- ...

Non-Parametric Regression



Parametric Regression

- Choosing **in advance** the **shape** of the function
- E.g., Linear regression

Non-Parametric Regression

- **No prior choice** on the **shape** of the function
- E.g., KNN-regression

Regression

KNN-regression

- An extension of KNN algorithm to **regression** case
- Simply compute target value as (weighted) **combination** of those of the K nearest neighbors

$$\hat{y}_j = \frac{\sum_{k \in \mathcal{N}_j} w_k y_k}{\sum_{k \in \mathcal{N}_j} w_k}$$

$$w_k = \frac{1}{d(x_j, x_k)}$$

Linear Regression



Linear regression with **least-squares** approach

Minimize the **Residual Sum of Squares** (RSS)

$$RSS = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

where \hat{y}_i is computed as

$$\hat{y}_i = \beta^T x = \sum_{i=1}^K \beta_i x_i$$

Linear Regression



In matrix form we are given N examples with K features: $Y \in \mathbb{R}^N, \beta \in \mathbb{R}^K, X \in \mathbb{R}^{N \times K}$

$$\min_{\beta} \|Y - \beta^T X\|^2$$

Linear Regression

The problem is **over-dimensional**: we have more equations than variables ($N \gg K$)...
There exist a closed-form solution

$$N \gg K$$

$$\beta^* = \boxed{(X^T X)^{-1} X^T} Y$$

MOORE-PENROSE
PSEUDO-INVERSE

Linear vs. Non-Linear


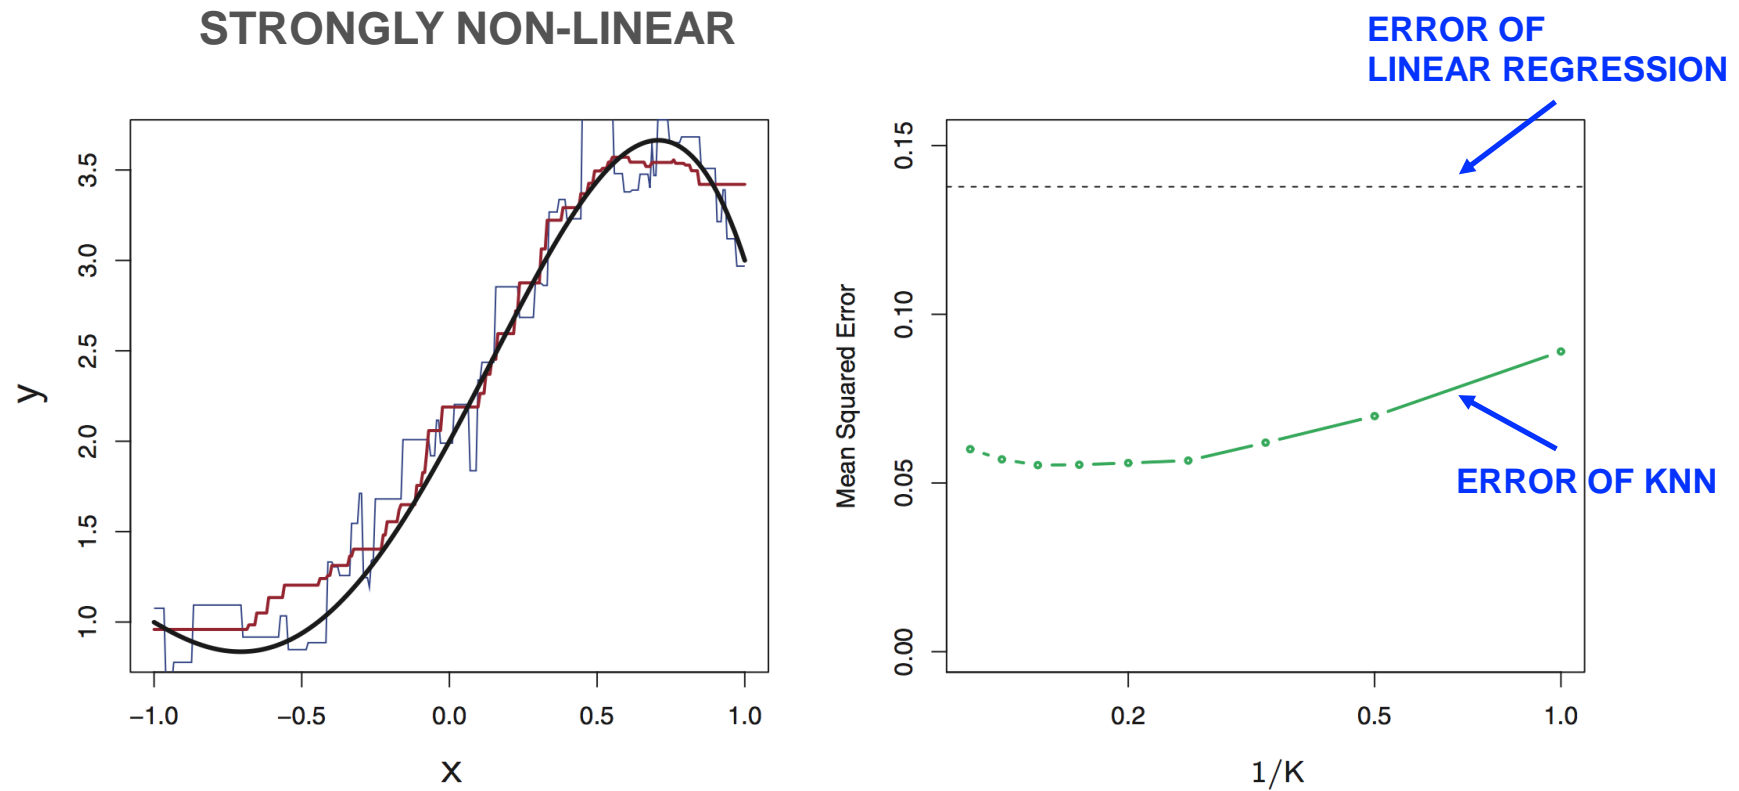


Figure from [James et al.]

Linear vs. Non-Linear



BLACK
True relationship

BLUE
KNN with $K=1$

RED
KNN with $K=9$

Figure from [James et al.]

More on Regression



Other regression models:

- Regression trees
- Neural networks
- ...

Regression Trees



An **extension** of decision trees to regression

- Use the **improvement** in Mean Squared Error to evaluate and select attributes during training
- In a leaf node, compute the **average** of the target variable across the training examples in the leaf
- Use **ensemble** techniques (Random Forests and Gradient Boosting) as for classification

Neural networks




Any neural architecture can be used for regression

- For Multi-Layer Perceptron, just use **one neuron** in the output layer, with linear activation function
- Use a **loss function** that evaluates regression tasks, such as Mean Squared (or Absolute) Error
- Specific architectures are dedicated to time-series analysis: i.e., **recurrent neural networks** such as Long Short-Term Memory (LSTM) networks

Example on Time-Series

Parameters to control when forecasting time-series?



Example with Scikit-Learn

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
```

```
df = pd.read_csv("yahoo_stock.csv", sep=",")
```

← **UNIVARIATE
TIME-SERIES**

```
data = df["Close"].values.astype('float32')
```

← **WINDOW = NUMBER OF FEATURES**

```
X = []
```

```
y = []
```

```
W = 10
```

```
for i in range(W, data.shape[0]):
```

```
    X.append(data[i-W:i])
```

```
    y.append(data[i])
```

```
X = np.array(X)
```

```
y = np.array(y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)
```

Example with Scikit-Learn

v0 v1 v2 v3 . . .

x0 x1 x2 x3 -> y = x5

x1 x2 x3 x4 -> y = x6

x2 x3 x4 x5 -> y = x7

. . .

. . .

Example with Scikit-Learn

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor

print("*** Random Forest ***")
clf = RandomForestRegressor()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print(mean_absolute_error(y_test, y_pred))

print("*** Linear Regression ***")
clf = LinearRegression()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print(mean_absolute_error(y_test, y_pred))

print("*** KNN Regression ***")
clf = KNeighborsRegressor(n_neighbors=3)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print(mean_absolute_error(y_test, y_pred))
```

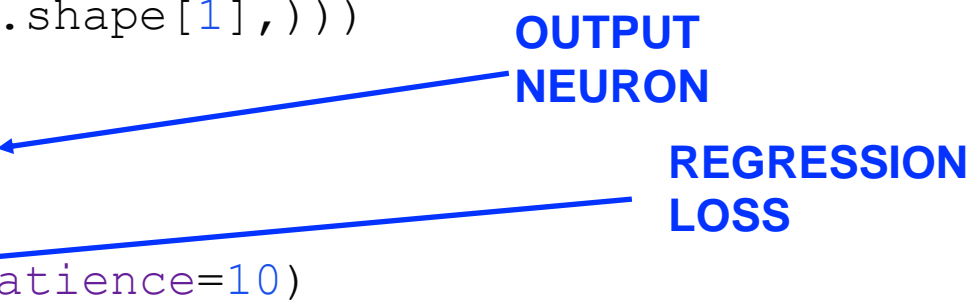
Example with Keras

```
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
```

```
model = Sequential()
model.add(Dense(50, input_shape=(X_train.shape[1],)))
model.add(Dense(20))
model.add(Dense(1, activation='linear'))
es = EarlyStopping(monitor='val_loss', patience=10)


model.compile(loss='mean_squared_error', optimizer='adam',
metrics=['mean_squared_error'])

model.fit(X_train, y_train, epochs=1000, batch_size=16,
validation_split=0.2, callbacks=[es])
```



OUTPUT NEURON

REGRESSION LOSS

- 
- <https://towardsdatascience.com/how-not-to-use-machine-learning-for-time-series-forecasting-avoiding-the-pitfalls-19f9d7adf424>